



[16.09.2021]

Bulky mediastinal lymphoma classification with ML-techniques

Activity status n. 1

Matteo Barbetti

University of Firenze
INFN - Firenze

Objective

Classification of *bulky mediastinal lymphoma* using **Machine Learning** techniques applied to PET and CT images.

Lymphoma types:

1. *Hodgkin (HL)*
2. *Gray Zone (GZ)*
3. *Primary Mediastinal Lymphoma (PML)*

PLANNED

From low-level data

- Data: PET and CT images
- Technique: Image classification
- Algorithm: *Convolutional Neural Nets*

IN PROGRESS

From high-level data

- Data: features drawn by LIFEx from images
- Technique: Binary and multiclass classification
- Algorithm: *Logistic regression + Random Forest*

First attempts: *high-level dataset*

Database of 119 rows x 101 columns

- **different data types** → need for homogeneous dataset
 - numbers
 - dates
 - texts
 - intervals
- **dim[instance space] ~ dim[feature space]** → need for reducing the feature space

Public repository on GitHub to **prepare**, **clean** and **study** the high-level dataset through Jupyter Notebooks:

- `data_preparation` -- data format correction
- `data_visualization` -- correlated features removal
- `binary_classification` -- classification in HL and non-HL
- `multiclass_classification` -- classification in HL, GZ and PML



[mbarbetti/LNHunter](https://github.com/mbarbetti/LNHunter)



data_preparation

Feature list

ID

Lymphoma typer (HL =1, GZ = 2, PML =3)

Data nascita

dPET staging

SUVmin (SUV)

SUVmean (SUV)

SUVstd (SUV)

SUVmax (SUV)

MTV (# vx)

MTV (mL)

SMTV (mL/Kg)

TLG (SUV*mL)

STLG (SUV*mL/Kg)

MTV (# vx) TOT

MTV (mL) TOT

SMTV (mL/Kg) TOT

TLG (SUV*mL) TOT

STLG (SUV*mL/Kg) TOT

SHAPE_Volume(mL)

SHAPE_Volume(vx)

SHAPE_Sphericity[onlyFor3DROI]

SHAPE_Surface(mm2)[onlyFor3DROI]

SHAPE_Compacity[onlyFor3DROI]

PARAMS_DistanceOfNeighbours

PARAMS_NumberOfGreyLevels

PARAMS_BinSize

PARAMS_IntensityResampling

PARAMS_BoundsRangeOfValueAfterDiscretization

PARAMS_ZSpatialResampling

PARAMS_YSpatialResampling

PARAMS_XSpatialResampling

NGLDM_Coarseness

NGLDM_Contrast

NGLDM_Busyness

GLZLM_SZE

GLZLM_LZE

GLZLM_LGZE

GLZLM_HGZE

GLZLM_SZLGE

GLZLM_SZHGE

GLZLM_LZLGE

GLZLM_LZLGE

GLZLM_GLNU

GLZLM_ZLNU

GLZLM_ZP

TimePosition

zLocation

CONVENTIONAL_SUVbwpeakSphere0,5mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CONVENTIONAL_SUVbwpeakSphere1mL

CHECK_Cluster(s)ToSmall

GLCM_Homogeneity[=InverseDifference]

GLCM_Energy[=AngularSecondMoment]

GLCM_Contrast[=Variance]

GLCM_Correlation

GLCM_Entropy_log10

GLCM_Entropy_log2[=JointEntropy]

GLCM_Dissimilarity

GLRLM_SRE

GLRLM_LRE

GLRLM_LGRE

GLRLM_HGRE

GLRLM_SRLGE

GLRLM_SRHGE

GLRLM_LRLGE

GLRLM_LRHGE

GLRLM_GLNU

GLRLM_RLNU

GLRLM_RP

- number
- date
- text
- interval

DISCRETIZED_SUVbwpeakSphere1mL(value only for PET or NM)

DISCRETIZED_TLG(mL)[onlyForPETorNM]

DISCRETIZED_HISTO_Skewness

DISCRETIZED_HISTO_Kurtosis

DISCRETIZED_HISTO_ExcessKurtosis

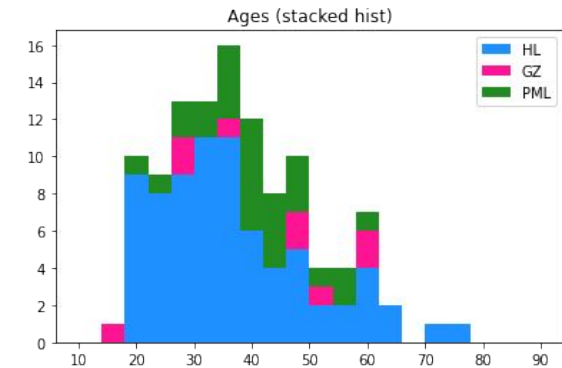
DISCRETIZED_HISTO_Entropy_log10

DISCRETIZED_HISTO_Entropy_log2

DISCRETIZED_HISTO_Energy[=Uniformity]

Dataset transformation

- CSV file → *Pandas DataFrame* (Python object)
- Features with no clear contents **dropped**
 - PARAMS_IntensityResampling
 - PARAMS_BoundsRangeOfValueAfterDiscretisation(SUVbw)
 - CHECK_Cluster(s)ToSmall
- **Data format correction**
 - decimal separator : “,” → “.”
 - date notation : “dd/mm/yyyy” → *datetime* (Python object)
- **New feature added**
 - **age**, from dates subtraction (then dropped)



99 → 95 FEATURES


data_visualization

Feature correlation [1/2]

Two features strongly **correlated** bring the same statistical information then, in order to **increase** the Machine Learning capabilities, one should **clean** the dataset to have independent features.


$$\rho_{X,Y} = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_j (x_j - \bar{x})^2 \sum_k (y_k - \bar{y})^2}}$$

Pearson correlation coefficient



Feature correlation [2/2]


- Pandas DataFrame is able to compute easily **pairwise correlation of columns** (features), excluding NA/null values
- Here a graphical example for a bunch of variables
 - MTV (# vx)
 - MTV (mL)
 - SMTV (mL/Kg)
 - MTV (# vx) TOT
 - MTV (mL) TOT
- Variables with strong correlations **dropped**
 - $\rho_{X,Y} \geq |0.75| \rightarrow$ strong correlation



95 → 17 FEATURES

binary_classification

Instance space



- Missing information
 - some features not available for all the instances
 - instances **dropped** to keep homogeneity
 - 119 → 101 rows (*instances*)
- Dataset clearly **unbalanced**
 - HL-class is over-represented w.r.t. the other two ones
 - classification suffers from unbalancing
- **Binary classification**
 - HL and non-HL classification
 - dataset a bit more balanced
 - NHL-HL ratio : 55.4%

Train/Test split

DATASET | size: 101 | ratio: 100%

TRAIN-SET | size: 80 | ratio: 80%

TEST-SET | size: 21 | ratio: 20%

To ensure that our Machine Learning model will be able to **generalize** well to new cases, namely to behave like a *true predictor*, we should avoid to use all the dataset to train our model. A better option is to **split the data into two sets**:

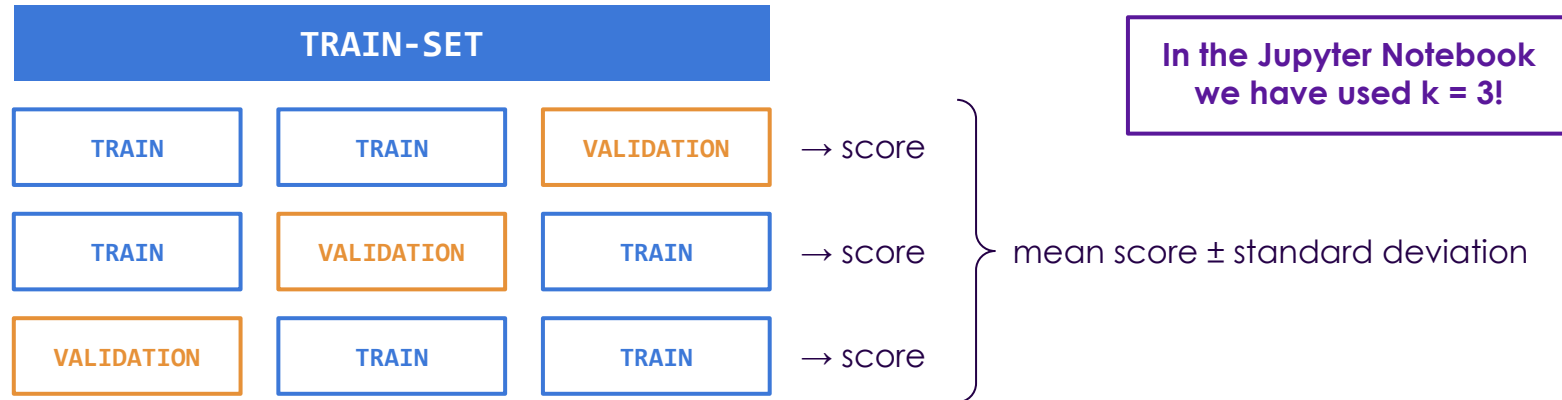
- **training set** to train the model
- **test set** to test the model

The error rate on new cases is called the **generalization error**, and we can estimate it evaluating the trained model on the test set.

Model optimization: *k*-Folds Cross Validation

In the search of the best possible predictor, we should avoid that our model learns by heart the training instances, because it will result in a **greater generalization error**.

A possible solution is to use ***k*-Folds Cross Validation**, an algorithm based on the splitting of the training set into k different subsets. We will use $k-1$ subsets to train the model and leave the last subset to validate it. Then, the average performance of our model against each of the folds can be used as a robust score to build an **optimization problem**, namely to find the best possible model.



Performance measures: accuracy, precision & recall

The **accuracy** is generally not the preferred performance measure for classifiers, especially when one are dealing with **unbalanced datasets**. In fact, classifiers succeed in recognizing instances belonging to most frequent label, resulting in high accuracy even if the other instances are mis-classified.

In order to ensure a model evaluation robust against unbalanced datasets, one can use **precision** or **recall** scores.

In the Jupyter Notebook
we have used RECALL!

		Predicted labels	
		Negative	Positive
Actual labels	Negative	True Negative (TN)	False Positive (FP)
	Positive	False Negative (FN)	True Positive (TP)

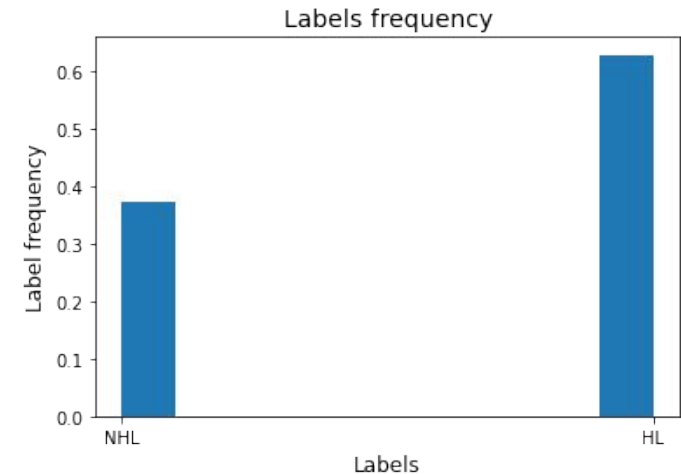
$$\text{accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{FP} + \text{TN} + \text{FN}}$$

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$


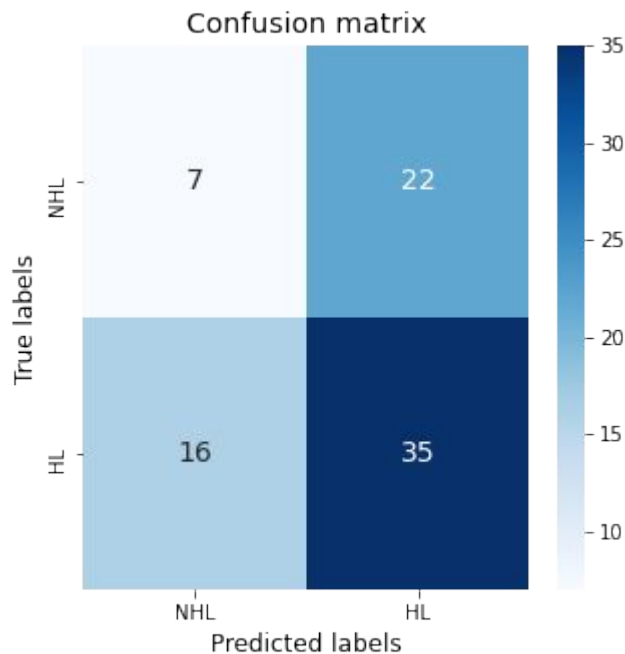
$$\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Dummy Classifier: *model baseline*


- The **Dummy Classifier** is a simple classifier that makes predictions using simple rules:
 - computes the frequency of each label
 - interprets such frequencies as probabilities
 - predicts new instances label on the basis of computed probabilities
- This classifier is useful as a simple baseline to compare with other (real) classifiers



Dummy Classifier: *performance*



Logistic regression



The **Logistic Regression** (LR) is a statistical model that in its basic form uses a **logistic function** $\sigma(\cdot)$ to model a binary dependent variable.

- LR model computes a weighted sum of the input features z (plus a bias term)
- This weighted sum are mapped it to a number between 0 and 1 through the logistic function
- The logistic output $\sigma(\mathbf{x})$ can be interpreted as the **probability** that \mathbf{x} belongs to the positive class

$$z(\mathbf{x}_i) = w_0 + \sum_{j=1}^n w_j x_{i,j} = \mathbf{x}_i^T \mathbf{w}$$


weighted sum of features

$$\hat{p}_i \doteq \sigma(z(\mathbf{x}_i)) = \frac{1}{1 + \exp(-z)}$$

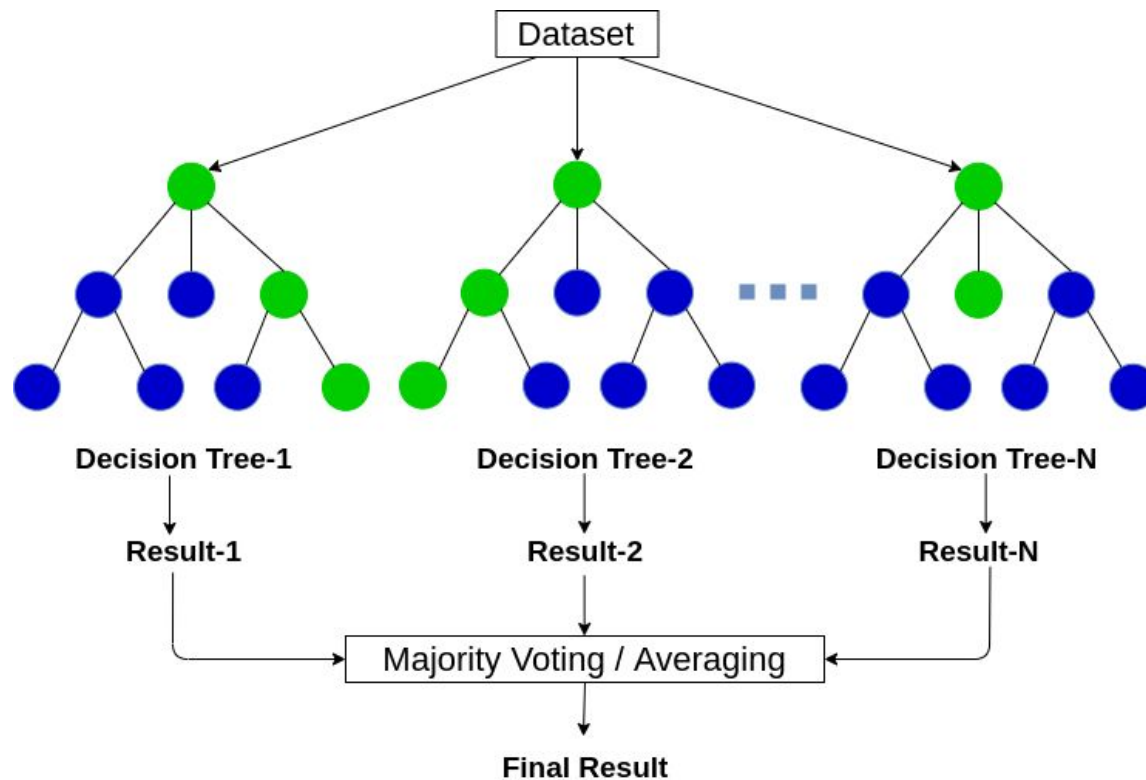
$< 0.5 \rightarrow \text{NHL}$
 $\geq 0.5 \rightarrow \text{HL}$

predictions


Logistic regression: *performance*



Random Forest classifier




Random Forest classifier: *performance*




Gradient BDT classifier

Single Decision Tree




Gradient Boosted Trees

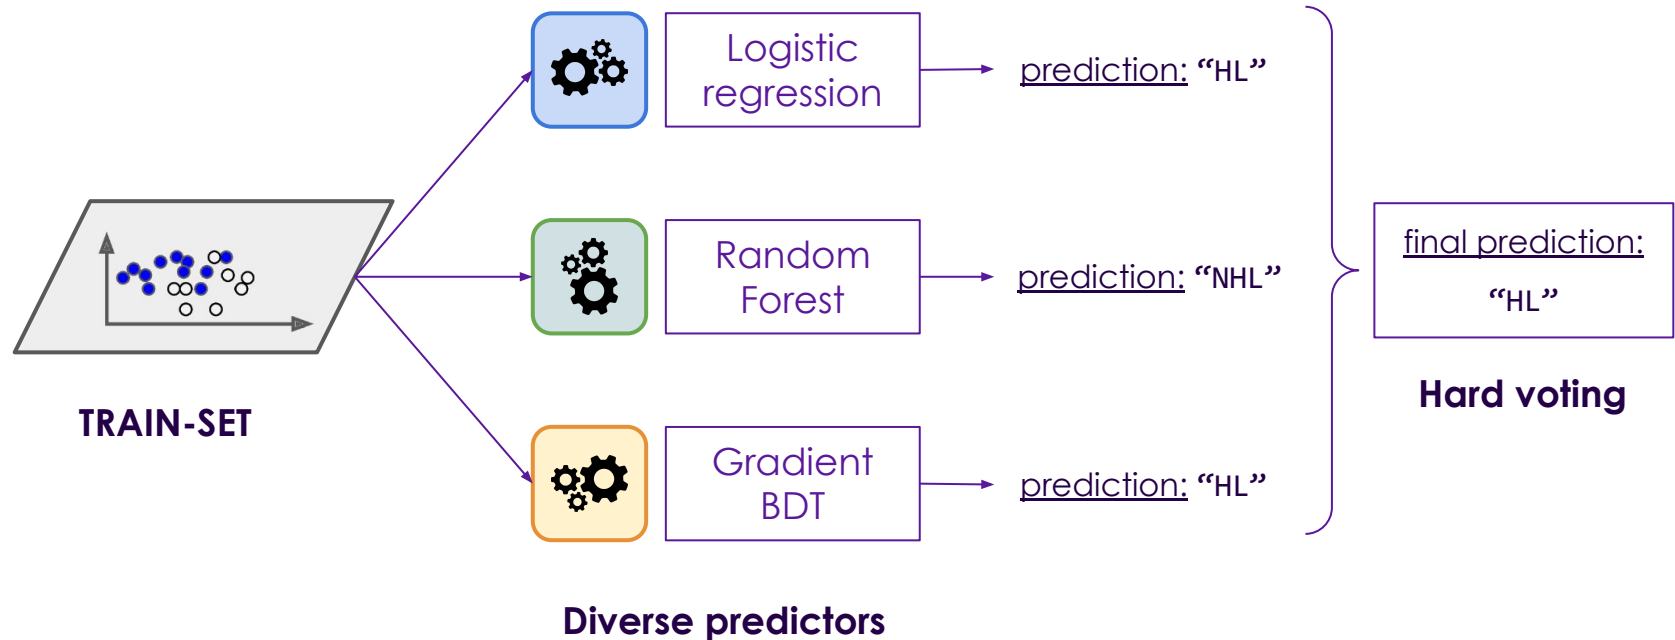


- **Gradient BDT** combines multiple simple decision trees into a stronger learner
- Each tree is trained on the **residuals** from the previous sequence of trees
- All trees are then combined together using an additive model


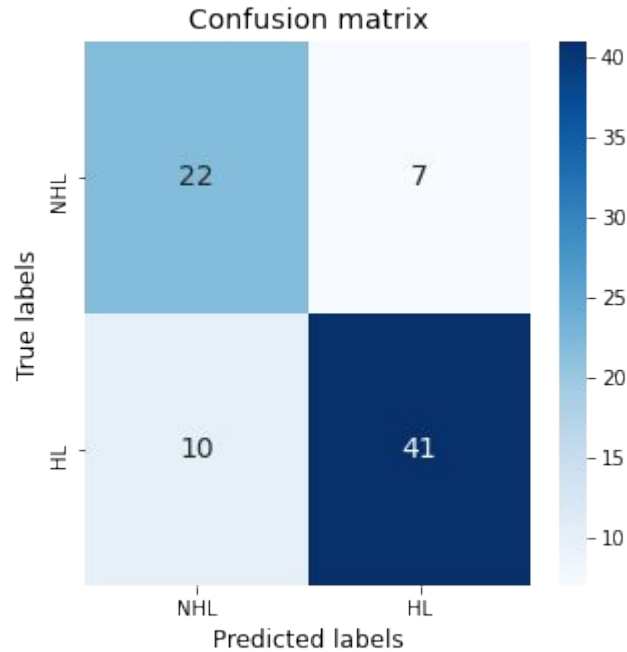
Gradient BDT classifier: *performance*





Model combination



Model combination: *performance*




Results on the test-set





multiclass_classification

Instance space





- Missing information
 - some features not available for all the instances
 - instances **dropped** to keep homogeneity
 - 119 → 101 rows (*instances*)
- Dataset strongly **unbalanced**
 - HL-class is over-represented w.r.t. the other two ones
 - classification suffers from unbalancing
- **Multiclass classification**
 - HL, GZ and PML classification
 - GZ-HL ratio : 13.8%
 - PML-HL ratio : 41.5%
- Train/Test split : **80%/20%**
- k-Folds Cross Validation with **k=3**


Dummy Classifier




Logistic regression: *performance*





Random Forest classifier: *performance*




Gradient BDT classifier: *performance*



Model combination: *performance*



Results on the test-set



Conclusion

- Preliminary studies
- There is room for improvements
 - **increasing** the dataset → more data
 - **balancing** the dataset → expanding *low-represented* classes
 - working with **low-level data** → *image classification* techniques
- Still need for taking a look at literature

Open to any kind of suggestions!

Backup

Undertraining & overtraining

